

Final Exam **with answers...** **...and comments**

CSYE7200 37076 Big-Data Sys Engr Using Scala SEC 01
—Spring 2016

Introduction—important

This exam is quite long. But it isn't really all that hard. I suggest you try to answer all the questions to the best of your ability and *only then* go back and check your results, assuming you have time. I don't want you to get stuck writing long essays or getting every little detail of code right before you go on to the next question. Attempting a question and getting in the "ball park" will get you many more points than leaving a complete blank (which earns a big fat zero).

In all your answers, I urge you to be brief! *Do not write essays*. Do not even write complete sentences unless that's the only way to be clear.

If you really need it, you have up until 9:00 pm (i.e. three hours). But I do hope you'll be finished before then!

The answers I have given are exactly as I wrote them up last evening during your exam with a few minor tweaks here and there to add clarity, formatting, or to give alternatives. Where the text is in bold, it implies that you should have **something equivalent to get full marks**. If it's not in bold but follows "or" then that's equivalent. Otherwise if you have something else which is correct, you might get points for extra credit.

Overall, I'm shocked by some of your answers. Apparently, quite a few of you really don't seem to understand Scala and/or functional programming at all.

I'm always surprised to find what's easy and what's hard. I thought that a lot of the questions would be really easy. For example, Questions 2, 7, 8, 10 and perhaps 6. The only one you all found easy was #6. #2 was, on the other hand, one of the least well answered yet the answer really is staring you in the face!

Prelude

One of the tricky things with Scala that most of us Java programmers find problematic, is how to rewrite algorithms that involve mutable state into functional programming (immutable) form.

Take for example this simple program:

```
object Count extends App {  
  val counter = Counter(10)  
  var i = 0  
  while ({i = counter.next; i>=0}) {println(i)}  
}  
case class Counter(var x: Int) {  
  def next = {x = x-1; x}  
}
```

There are two mutable objects here: *counter* and *i*. The two require slightly different techniques to refactor into a FP form.

Question 1

1. Given a (generator) class $A[T]$ with *mutable* state and method *next*: T , show, in general terms, how to refactor it into an *immutable* class $B[U]$ that also has a method which you can use to get its next U value. Your solution must be referentially transparent. Hint: think back to Week 2 of the class.
2. In general terms, show how to refactor an iterative process using a *var* x : Int into a referentially transparent (FP-style) process. Hint: think back to Week 4 (quick review, part 2)
3. Using both (preferably) of the techniques you describe in 1 and 2, rewrite *Counter* and the *Count* object in a referentially transparent way.

1. (6) **The *next* method must return both the next U value and the next $B[U]$ value** (which will be used to give us the following U value); Alternatively, you can return just the next B value provided that B as a method to access its U value. Some of you took a rather literal view of my hint: including stuff about random number generators (seeds, etc.): you must be able to adapt from solutions to other, similar problems, not clone them!
2. (8) **we need to create a (private) *tail-recursive* method that will include among its parameters both x and whatever is used to terminate the loop.** Any cumulative results will be collected in an immutable collection which will typically be a third parameter. Some of you came up with an interesting approach: make an immutable list of all the values that could arise from calling $B.next$ and iterate through those. But that will only work if you happen to know in advance the sequence that B is going to yield! What if B is a random number generator?
3. (11) **case**

```
class B[U](u: U)(implicit nextU: U=> U) {  
  def next = {val v = nextU(u); (v, B(v)(nextU))}  
}  
def loop(b: B, f: (Int)=>Unit): B = b.next match {  
  case (0,b2) => b2  
  case (x2,b2) => f(x2); loop(b2,f)  
}  
loop(B(10),println(_))
```

[Note that I put in a mechanism to get the next value of u in case U is not an Int . It's a generalization that I didn't expect you to put in — and in any case it could be done in several different ways.]

Assignment 1: Document

You have been asked to work on a data structure that was close to being completed by a colleague who suffered an unfortunate accident. The code that you have been given looks like this:

```
package edu.neu.coe.scala
trait Document[K, +V] extends (Seq[K] => V) {
  def get(x: Seq[K]): Option[V]
}
case class Leaf[V](value: V) extends Document[Any, V] {
  def get(x: Seq[Any]): Option[V] = x match {
    case Nil => Some(value)
    case _ => None
  }
}
case class Clade[K, V](branches: Map[K, Document[K, V]]) extends
Document[K, V] {
  def get(x: Seq[K]): Option[V] = x match {
    case h::t => branches.get(h) flatMap {_.get(t)}
    case Nil => None
  }
}
```

When you try to compile it you see essentially the same error for both *Leaf* and *Clade*:

class Leaf needs to be abstract, since method apply in trait Function1 of type (v1: Seq[Any])V is not defined.

Question 2

1. Explain succinctly why you are seeing this error.
2. Fix the error in the most simple, obvious and elegant manner which is also consistent with the *scala.collection.immutable.Map* trait.

1. (4) Because *Document* extends *Function1[Seq[K], V]* we need a *def apply(ks: Seq[K]): V* method.

2. (8) *def apply(ks: Seq[K]): V = get(ks).getOrElse(default())*
def default() = throw new NoSuchElementException

=====

Part of the Scalatest specification looks like this:

```
class DocumentSpec extends FlatSpec with Matchers {
  "Leaf(1).get" should "yield Some(1) for Nil, None for i" in {
    val doc = Leaf[Int](1)
    doc.get(Nil) should matchPattern { case Some(1) => }
    doc.get(Seq("i")) should matchPattern { case None => }
  }
```

```

}
"Clade.get" should "work appropriately for one level" in {
  val one = Leaf[Int](1)
  val doc = Clade(Map("one" -> one))
  doc.get(Nil) should matchPattern { case None => }
  doc.get(Seq("one")) should matchPattern { case Some(1) => }
}
it should "work appropriately for two levels" in {
  val one = Leaf[Int](1)
  val doc1 = Clade(Map("one" -> one))
  val doc2 = Clade(Map("a" -> doc1))
  doc2.get(Nil) should matchPattern { case None => }
  doc2.get(Seq("a", "one")) should matchPattern { case Some(1) => }
}
}

```

Unfortunately, although the *Document* module itself is now compiling, the spec file does not. In particular, the expression *Map("one" -> one)* shows the following error:

Multiple markers at this line:

- type mismatch; found :
scala.collection.immutable.Map[String,edu.neu.coe.scala.Leaf[Int]] required:
Map[Any,edu.neu.coe.scala.Document[Any,Int]] Note: String <: Any, but trait Map is
invariant in type A. You may wish to investigate a wildcard type such as `_ <: Any`.
(SLS 3.2.10)
- etc.....

Question 3

1. Explain succinctly why you are seeing this error.
 2. Fix the error in the most simple, obvious and elegant manner (incidentally, this will not be consistent with the *scala.collection.immutable.Map* trait).
- 1. (6) The *doc* variable is inferred to be a *Seq[String]=>V* (because of key: "one") but *one* is a *Seq[Any]=>Int*. *Any* is a super-type of *String* (because this is a contra-variant position), but *Seq[Any]* can only be a super-type of *Seq[String]* if *K* is contravariant. I reworked this answer to try to make it clearer. It only goes to show that this question was really too hard!**
- 2. (3) we add a "-" in front of *K* (after *Document*)**

Question 4

1. Explain the purpose of this data structure.

2. Describe exactly what is going on in the right-hand-side of the case (in *Clade*) which matches the pattern $h::t$. For what purpose is *flatMap* used here?
3. Suppose you were asked to implement a method in the *Document* trait which has the following signature:

```
def get(x: String): Option[V]
```

such that x is split into Strings delimited by “.”. For example, $get("a.b")$ would be the equivalent of $get(List("a", "b"))$.

Describe the components you will need to make this work (I do *not* need actual code). Don’t forget that *Document* is defined for keys of type *Seq[K]*—not *Seq[String]*.

1. (5) It provides us a hierarchical key-value store of unlimited depth;
2. (5) the result of *branches.get(h)* is an *Option[Document[K,V]]*. *FlatMap* allows us to (recursively) invoke *get(t)* on an actual document (i.e. a *Some*) otherwise simply return *None* if h doesn't match a key;
3. (9) we need:
 - A. a parser to split the string up (suggest *StringOps.split("""\.""")*)
 - B. a means of turning an array into a *Seq* (suggest *toList.toSeq*)
 - C. an *implicit val String=>K* which will be (implicitly) invoked on a *String* to provide a *K*

=====

Intermezzo: Spark

Question 5

1. Name at least two of the most significant characteristics of *RDDs*.
 2. Explain the differences between the two types of operations on an *RDD*: actions and transformations.
 3. Give one example of an action and one example of a transformation.
 4. Given that, in general, *RDDs* have dependencies on other *RDDs*, how do you think transformations on *RDDs* are implemented. Put your thoughts concisely in words (not code), using a simple method as an example.
-
1. (5) **Lazy, parallelizable** (or partitioned/distributed), immutable, monad, etc. etc.;
 2. (4) **actions turn the *RDD* into a strict value; transformations create a new *RDD* based on the original *RDD***; My sense is that many of you don't really understand that there is a fundamental difference here. It's best illustrated in the Spark REPL: do anything with a bunch of transformations (or invoke parallelize) and the REPL will respond immediately. Do an action and the underlying map/reduce (Yarn, Mesos, whatever) will spring into action and you will have to wait several seconds depending on the complexity. Even `sc.textFile("non-existent file")` will appear to work until you invoke an action.
 3. (3)
 - A. **collect** or any of... *reduce, count, first, take, takeSample, takeOrdered, saveAsTextFile, saveAsSequenceFile, saveAsObjectFile, countByKey, foreach*;
 - B. **map** or any of... *flatMap, filter, mapPartitions, mapPartitionsWithIndex, sample, union, intersection, distinct, groupByKey, reduceByKey, aggregateByKey, sortByKey, join, cogroup, cartesian, pipe, coalesce, repartition, repartitionAndSort...* etc.;
 4. (6) **for the *map* method on *RDD[T]* which takes as a parameter *f: T=>U*, we create a new *RDD* which simply references the original *RDD* and remembers that we want to apply *f* to the elements (without actually applying *f*)**. If we apply another function *g* we will compose the functions together using *f andThen g* and remember that function instead (the same as we did with *LazyNumber* in assignment 5).

Question 6

Write a short Spark program to count the number of words in a text file, excluding the following "stop" words: *the, a, an, in, of, by, and, at, by, with*. Assume that you are given a *SparkContext* (as you would be in Zeppelin, for instance). Pseudo-code is acceptable.

(13)

```
val r = sc.textFile("/Users/scalaprof/flatland.txt")
```

```
val stopWords = List("the", "a", "an", "in", "of", "by", "and", "at",
"by", "with")

val s = r flatMap {w => w.split("""\p{Punct}?\s+""")} filter (w => !
stopWords.contains(w.toLowerCase))

println(s.count)
```

I'm curious why so many did a key-based word count which is nice but not actually requested.

Assignment 2: Bridge Movement

[Remember, you only need to read this to the extent that you need clarification]

Contract bridge is a card game where two pairs of players compete, sitting around a table. Each player is dealt 13 cards and he and his partner try to win “tricks” (each made up of four cards). Bridge can be played for money or for fun. Duplicate bridge is a highly competitive card game in which players play the exact same deals as other players. So that this can be achieved we place the cards after play into a holder with four pockets—one for each player’s cards—these holders are called “boards”. If no boards or players ever moved, it wouldn’t be duplicate, so at least the boards always move. The plan for how boards or players move is called the “movement.” The two most common forms are “pairs” and “teams” but in the normal team form, only the boards move so there is no complexity.

You have been asked to write a Scala program to show how the movement works in a *pairs* game. The pairs are divided (somewhat arbitrarily) into two groups according to which direction they sit: “N/S” (north/south) and “E/W” (east/west). So, a bridge “encounter” involves four “things”: the N/S and E/W pairs, the set of board(s), i.e. deals, which they play and the table at which the play occurs.

You therefore have four corresponding quantities:

- t : the table number ($0..T-1$)
- n : the N/S pair number ($0..N-1$)
- e : the E/W pair number ($0..E-1$)
- b : the number of the set of board(s) to be played. ($0..B-1$)

A game is divided (in the time dimension) into “rounds”. For any given round, there will be an encounter at each of the T tables. After each encounter, three of these things move by a pre-determined step. The tables (the least movable objects) stay where they are. So, we have the following three quantities:

- dN : the number of tables (in a positive direction) which N/S will move
- dE : the number of tables (in a positive direction) which E/W will move
- dB : the number of tables (in a positive direction) which the boards will move

By convention, dB is always -1 in duplicate bridge, but we won’t assume that because we want to write something very general. And in most movements, $dN=0$ (the N/S pairs are stationary). The “perfect” bridge session has the following properties: $T = N = E = B = 13$; $dN = 0$, $dE = 1$, $dB = -1$. In this setup, there are two boards in each set, so players play 26 hands of bridge during about 200 minutes of play (this allows for 15 minute rounds, with a short break about halfway).

However, your task is to list the encounters that occur in each round of a more general game. In particular, there is a kind of movement which involves “phantom” tables where the encounters

do not count. This is called a “Howell” movement after the Massachusetts native who developed it—without the help of computers—about 100 years ago.

In this movement, $dN = -3$, $dE = -2$, $dB = -1$. If these numbers are co-prime with T then all proceeds smoothly. However, in the case where $T=9$, the N/S pairs would return to their starting table after three rounds. This requires something different in the movement such as setting $dN = -2$ or -4 after three rounds. Depending on how we do this, we might now reencounter an E/W pair from earlier so we might need to change the E/W movement after some number of rounds.

Since we are typically operating on things in threes (pertaining to N/S, E/W, and board respectively) we have chosen to create a type called a *Triple* which extends *Product3*:

```
case class Triple[T](_1: T, _2: T, _3: T) extends Product3[T,T,T] {  
  def map[U](f: T=>U): Triple[U] = ???  
  override def toString = s"n:${_1} e:${_2} b:${_3}"  
}
```

Question 7

1. Explain the key differences between a *TupleN* and a *ProductN* (you will most likely want to look them up in the Scala API to answer this question). Hint: think about why we didn't define *Triple* to extend *Tuple3*.

2. Why are the names of the three identifiers in *Triple* chosen as “_1”, etc.? What would happen if we simply call them “ns”, “ew”, “bd”?

3. Implement *map* on *Triple*:

```
def map[U](f: T=>U): Triple[U] = ???
```

4. Is *Triple* a monad? Justify your answer.

5. Given, the definition of *Triple* above, would you expect the following code to compile and work?:

```
def toStrings = for (p <- this) yield p.toString
```

6. What about this code? Explain why it does or does not compile:

```
def toStrings = for (p <- this; n <- this) yield p.toString
```

1. (4) ***TupleN* is a case class** (with all that implies, including no sub-classes) **which extends *ProductN*** and mixes in *Product* with *Serializable*. *TupleN* has explicit parameter/field names of *_1*, *_2*, ... **On the other hand, *ProductN* is a trait** with only a few concrete members such as *productArity*, and which is extended (synthetically) by all case classes— Many of you talked about how *Product* represents a Cartesian product—yes, this is the origin of the name but it isn't really very relevant to the question asked. You further described these things as if they were completely different: yet you know (or should) that *TupleN* extends *ProductN*!

2. (2) **because they would be undefined otherwise (they are defined as abstract in *Product3*);**

3. (5) **def map[U](f: T=>U): Triple[U] = Triple(f(_1), f(_2), f(_3))**

4. (2) **no—it lacks *flatMap* and *filter*;**

5. (2) **yes—it implements *map***; Many of you thought that this would not work. My answer to you is: **E pur si muove** (look it up).
6. (2) **no—it doesn't implement *flatMap***.
- =====

Question 8

1. Why do we need the “type” keyword in Scala?
 2. What use is it if we can only define it within a package or class? i.e. what do we need to do to bring it into scope?
1. (5) **For three reasons: as an alias for a type which is a non-simple type**; or to indicate the type of a Scala (singleton) “object”; to define a type *member* of a trait (or abstract class)—an alternative way of making classes polymorphic, as opposed to parametric polymorphism (i.e. generics). [Note that I forgot the last of these when I first wrote up the answers—we haven’t actually talked about this in class because it’s not really necessary unless you are going to design Scala libraries.]
2. (3) **we can import the definition of the type into our own scope**. Alternatively, we can create a “package object” [which we didn’t talk about much in class] and put the type definitions in there.
- =====

We model the moves using a *Stream[Trio]* where *Trio* is defined thus:

```
type Trio = Triple[Int]
```

The source of these moves is a set of three *Seq[Int]* objects, one for each of the three things to be moved. For our situation, we have:

```
val nsMoves = Seq(-3, -3, -2)
val ewMoves = Seq(-2)
val bdMoves = Seq(-1)
```

There are two additional types defined:

```
type MovePlan = Triple[Seq[Int]]
type Moves = Triple[Stream[Int]]
```

Question 9

In order to make use of these, we need to write some methods for the *Triple* object:

1. implement the following method to convert a *Triple* of *Streams* into a *Stream* of *Triples*:

```
def zip[U](ust: Triple[Stream[U]]): Stream[Tuple[U]] = ???
```
2. Implement the following method (presumably invoking your new *map* method) to use the elements each *Seq* provided to populate the corresponding *Stream*:

```
def toStreams[U](ust: Triple[Seq[U]]): Triple[Stream[U]] = ???
```

1. (10) `def zip[U](ust: Triple[Stream[U]]): Stream[Triple[U]] = for
 (((x,y),z) <- (ust._1 zip ust._2 zip ust._3)) yield Triple(x,y,z) [or
map alternative]. Another way very elegant way to write it is this (bonus point if you
 got this): def zip[U](ust: Triple[Stream[U]]): Stream[Triple[U]] = ust
 match {
 case Triple(a,b,c) => for {
 aa <- a
 bb <- b
 cc <- c
 }yield (Triple(aa, bb, cc))
 }`
2. (10) `def toStreams[U](ust: Triple[Seq[U]]): Triple[Stream[U]] = ust
 map { case us=>Stream.continually(us).flatten }`

=====

Another class we will need is *Encounter* defined thus:

```
case class Encounter(table: Int, n: Int, e: Int, b: Int)(implicit tables:
Int) {
  def move(moves: Trio, current: Position): Encounter = {
    val x = moves map {i => current.encounters(modulo(table-i)-1)}
    Encounter.fromPrevious(table,x)
  }
  override def toString = s"T$table: $n-$e@#$b"
}
```

Question 10

1. In order to figure out the moves, we need to add a method *modulo* which transforms a number *n* in the range *1-T...infinity* into the range *1..T*. Implement the body of this method (note that we start counting all of our objects that move from *one*, not *zero*).
2. You may not have seen an implicit parameter as part of a case class before but recall that the case class definition essentially just specifies how the constructor looks. Everything else is inferred from that. Why do you think the designer chose to make *tables* implicit while the other four parameters are explicit?
3. How does this kind of implicit value get satisfied?

1. (6) `def modulo(n: Int) = (n-1+tables) % tables + 1`
2. (3) **because in an application, *tables* doesn't change so we can set up a starting set of *Encounter* objects economically (i.e. without having to specify *tables* for each one).**
3. (4) **it gets satisfied by a statement in scope of the form *implicit val tables = T***

=====

After the movement (and play) is complete, we need to score the event to see who won. Each board (traditionally) carries with it a “traveler”, a slip of paper that records the scores, pair numbers, etc. at each table. I say traditionally because we usually using an electronic scoring device nowadays. Imagine that we could scan the travelers at the end of the session and input them using a bridge-DSL. Here’s (part of) the code we need (“recap” is the term used for all of the travelers collected together):

```
class RecapParser extends JavaTokenParsers {
  // XXX traveler parser yields a Traveler object and must start with a "T"
  // and end with a period. In between is a list of Play objects
  def traveler: Parser[Traveler] = "T"~>wholeNumber~rep(play)<~"."~"" ^^
  {case b~r => Traveler(Try(b.toInt),r)}
  // XXX play parser yields a Play object and must be two integer numbers
  // followed by a result
  def play: Parser[Play] = wholeNumber~wholeNumber~result ^^ { case n~e~r =>
  Play(Try(n.toInt),Try(e.toInt),r) }
  // XXX result parser yields a PlayResult object and must be either a number
  // (a bridge score) or a string such as DNP or A+-
  def result: Parser[PlayResult] = (wholeNumber | "DNP" | regex("A[-\\
+]?[0-9].r) ) ^^ { case s => PlayResult(s) }
}
```

Question 11

For these questions, you might find the following link helpful: <http://www.scala-lang.org/files/archive/api/2.11.2/scala-parser-combinators/>. In your answers, I want *key points*, not essays!

1. Explain briefly what “~” signifies. Your explanation should take into account the fact that we see it on both sides of the “^^”.
 2. Explain briefly what “~>” signifies.
 3. Explain briefly what “^^” signifies. Include the type of *wholeNumber*, for example, as part of your explanation.
-
1. (7) ~ is a case class with two parameters of type *Parser[T]* (the left side and right side). It successfully parses if the two parsers match in sequence. **On the right hand side of ^^, ~ is used in a pattern-matching situation** (i.e. its *unapply* method is being invoked). Many of you ignored (or missed) the bit about appearing on both sides of “^^”. The only way that it could, syntactically, appear in both places is if it’s a case class (or any class with an extractor).
 2. (4) ~> is an operator (method) on *Parser[T]* which, like ~, must satisfy both operands in sequence but results in just the right-hand parameter (i.e. it only captures the right side).
 3. (6) ^^ is an operator on *Parser[T]* which transforms it into a *Parser[U]*. The right hand side is a function of form *T=>U*. *wholeNumber* is a *Parser[String]*. ^^ behaves very much like *map* does for a functor (or monad). In fact, it turns out that ^^ actually invokes *map* which is in fact defined for *Parser[T]* (as is *flatMap*). I

don't think anyone satisfactorily mentioned the transformation of *wholeNumber* as an example

=====

The way duplicate is scored is that, for any one board, each pair gets two points for every pair they beat and one point for every pair that they tie. So, let's say a board is played 13 times. One pair achieves a better score than all of the other pairs sitting in their direction. They get 24 "matchpoints". This is called a "top." Let's say all of the twelve other pairs get the same result as each other. They each score 11 matchpoints. The total matchpoints for the board adds up to $24 + 12 * 11$ which equals 156. This is correct (top * pairs).

The *Traveler* class is defined thus:

```
case class Traveler(board: Int, ps: Seq[Play]) {  
  def isPlayed = ps.nonEmpty  
  ...  
}
```

In the *Traveler* class, there is a method to matchpoint a play:

```
def matchpoint(x: Play): Option[Rational] = if (isPlayed) {  
  val isIs = (for (p <- ps; if p != x; io = p.compare(x.result); i <-  
io) yield (i,2)) unzip;  
  Some(Rational.normalize(isIs._1.sum, isIs._2.sum))  
}  
else None
```

where *Play* is defined thus:

```
case class Play(ns: Int, ew: Int, result: PlayResult) {  
  def compare(x: PlayResult): Option[Int] = ??? // if Int, then 0, 1, or 2  
  def matchpoints(t: Traveler): Option[Rational] = ???  
}
```

Question 12

1. What is the type of *isIs*?
2. Explain clearly how *isIs* is evaluated, describing what each component of the expression signifies.
3. Why is there a ";" at the end of the definition? What would happen if we removed it?
4. Can you think of *one* simple way to avoid having ";"? I can think of three, maybe four, ways, one of which is not simple.

1. (5) Type of *isIs* is *(Seq[Int], Seq[Int])*

2. (8) it is evaluated in two parts:

A. a for-comprehension that generates *p*: *Play* from *ps*; checks that *p* doesn't equal *x*; invokes *p.compare(x.result)* which yields *io* (an *Option[Int]*); from this we generate *i*, an *Int* (if it is a *Some* else we ignore it); finally yielding a *Tuple2* of *i* and 2.

- B. the result of the first part is a `Seq[(Int,Int)]` since `ps`, the first generator, is a `Seq[Play]` and this is unzipped into a `(Seq[Int],Seq[Int])`.**
3. (4) **because `unzip` actually takes a parameter (it's implicit) so the compiler thinks the next line might be specifying that parameter.** It thus thinks the whole expression is self-recursive.
4. (3) **put a `".`" before `unzip`;** put an extra blank line after the `val is/s` line; provide an actual implicit value; `import postfixOps` (?).