

3.3 Parallel Processing & Mutable State

© 2019 Robin Hillyard



Northeastern
University

Parallel Processing (1)

- Modern computers have multiple “cores” and are typically arrayed in clusters. This allows us (in practice, forces us) to do things *in parallel*. There are two types of parallel processing and *Scala* can be effective in both:
 - *immutable state*—the definition of the problem is known at the start and does not change until the problem is solved
 - e.g. count the total number of words in 1,000,000 documents
 - Amdahl’s law applies to this type of problem:
 - maximum possible speedup $\sigma = 1/\alpha$ where α is the proportion of total time which is non-parallelizable
 - *Because each document is independent of the others*, a master node can divide up (this part is not parallelizable) the documents to be processed by a set of worker nodes/threads;
 - Then (in parallel) each worker counts the words in the documents it was given and returns the result, asynchronously of course, as a future value
 - Once all of these future values are realized, they can be summed to get the grand total (that final step is also non-parallelizable)
 - This is (more or less) the principle on which Map/Reduce works

Parallel Processing (2)

- continuing:
 - *mutable state*—the conditions are inherently changing at all times
 - e.g. ticket agency
 - there are two fundamental mutable states: the pool of tickets and the bank account of agency
 - it is therefore essential that these can be updated only by one thread at a time
 - there are other temporary mutable states—e.g. the status of an shopping cart—but these, if lost or corrupted, can be restarted
 - for these applications, we use *actors* (more about this later)


Parallel Processing (3)

- Apart from the internals of actors, is there any need to use mutable variables (vars) or mutable containers/collections?
- Caching (memoization)?
 - Maybe. But this can also be done inside an actor.
- Aggregation?
 - Normally, in functional programming we perform aggregation by recursion—no mutable state required.
- Sorting as a prelude to searching?
 - Yes—we do typically perform sorting on a mutable array but this again can/should be done inside an actor.
- Non-deterministic algorithms?
 - Yes, like the Newton Approximation or maybe genetic algorithms, etc.

Immutable list



Mutable list



Oops!

But surely we need *some* variables?

- Rarely! Although Scala provides *var* and has a library of mutable collections—you typically don't need them!
- Let's think about a box office app (see next slide):
 - if we are careful we can combine all mutable state into one *var* (the “state” of an actor);
 - when we take a list of items from another list (for example, we remove some tickets from the pool), surely we end up with lots of redundant copies of tickets! No, we don't.
 - in fact, the opposite is true: in an imperative program (e.g. Java) we typically end up doing a lot of list copying because our lists are mutable when they don't need to be!

Ticket Agency

```
package edu.neu.coe.scala
package boxOffice
object BoxOffice {
  val initialState = State(List(Ticket(1,1,100),Ticket(1,2,100)), List())
  def main(args: Array[String]): Unit = Tickets(initialState).start
}
case class Tickets(private var state: State) extends scala.actors.Actor {
  def act() {
    while (state.availability) {
      receive {
        case sale: Sale => state = state.makeTransaction(sale)
        case Status => sender ! state
        case _ => println("unknown message")
      }
    }
  }
}
case class State(tickets: List[Ticket], transactions: List[Transaction]) {
  def availability = tickets.length > 0
  def makeTransaction(sale: Sale): State = ???
}
case class Ticket(row: Int, seat: Int, price: Int)
case class Sale(tickets: List[Ticket], transaction: Transaction)
case class Transaction(creditCard: Long, total: Int, timestamp: Long, confirmation: String)
case object Status
```

??? is a boon to designers
who want to create stubs
to be implemented later

Scala actors have been
deprecated in favor of Akka

The field(s) of a case class can be marked "var"

Take care of mutable state when testing

```
package edu.neu.coe.scala;
import static org.junit.Assert.assertEquals;
import org.junit.BeforeClass;
import org.junit.Test;
import scala.util.Random;

public class GenericTest {
    private static Random random;
    @BeforeClass
    public static void setUpBeforeClass() throws Exception {
        random = new Random(0L);
    }
    @Test
    public void test1() {
        int x = 0;
        for (int i = 0; i < 2; i++)
            x = random.nextInt();
        assertEquals(x, -723955400);
    }
    @Test
    public void test2() {
        int x = random.nextInt();
        assertEquals(x, 1033096058);
    }
}
```

This is Java, by the way



what if we change this number? How will test2 work out?



what if the test runner decides to run test2 before test1?



- We will come back to this problem later.